

ML Assignment 1

Perform data preprocessing and exploratory data analysis on a given dataset to prepare it for machine learning modeling.

```
In [21]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Load CSV

```
In [22]: df = pd.read_csv("D:/Machine Learning/titanic - titanic.csv")
df.head()
```

Out[22]:

	Survived	Pclass	Name	Sex	Age	Siblings/Spouses Aboard	Parents/Children Aboard	Fare
0	0	3	Mr. Owen Harris Braund	male	22.0	1	0	7.2500
1	1	1	Mrs. John Bradley (Florence Briggs Thayer) Cum...	female	38.0	1	0	71.2833
2	1	3	Miss. Laina Heikkinen	female	26.0	0	0	7.9250
3	1	1	Mrs. Jacques Heath (Lily May Peel) Futrelle	female	35.0	1	0	53.1000
4	0	3	Mr. William Henry Allen	male	35.0	0	0	8.0500

Data Understanding

```
In [23]: df.shape
```

Out[23]: (887, 8)

In [24]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 887 entries, 0 to 886
Data columns (total 8 columns):
 #   Column      Non-Null Count Dtype  
 --- 
 0   Survived    887 non-null   int64  
 1   Pclass       887 non-null   int64  
 2   Name         887 non-null   object  
 3   Sex          887 non-null   object  
 4   Age          887 non-null   float64 
 5   Siblings/Spouses Aboard 887 non-null int64  
 6   Parents/Children Aboard 887 non-null   int64  
 7   Fare         887 non-null   float64 
dtypes: float64(2), int64(4), object(2)
memory usage: 55.6+ KB
```

In [25]: df.describe()

Out[25]:

	Survived	Pclass	Age	Siblings/Spouses Aboard	Parents/Children Aboard	Fare
count	887.000000	887.000000	887.000000	887.000000	887.000000	887.000000
mean	0.385569	2.305524	29.471443	0.525366	0.383315	32.30542
std	0.487004	0.836662	14.121908	1.104669	0.807466	49.78204
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.00000
25%	0.000000	2.000000	20.250000	0.000000	0.000000	7.92500
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.45420
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.13750
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.32920

In [26]: *# Check null values*
df.isnull().sum()


Out[26]:

```
Survived      0
Pclass        0
Name          0
Sex           0
Age           0
Siblings/Spouses Aboard 0
Parents/Children Aboard 0
Fare          0
dtype: int64
```

Exploratory Data Analysis


1. Target value distribution

```
In [27]: sns.countplot(x='Survived', data=df)
plt.title("Survival Count")
plt.show()
```




2. Survival VS Gender

```
In [28]: sns.countplot(x='Sex', hue='Survived', data=df)
plt.title("Survival based on Gender")
plt.show()
```




3. Survival VS Passenger Class

```
In [29]: sns.countplot(x='Pclass', hue='Survived', data=df)
plt.title("Survival based on Passenger Class")
plt.show()
```




4. Age Distribution

```
In [30]: sns.histplot(df['Age'], bins=30,kde=True)
plt.title("Age Distribution")
plt.show()
```



5. Fare VS Survival

```
In [33]: sns.boxplot(x='Survived', y='Fare', data=df)
plt.title("Fare VS Survival")
plt.show()
```



Data Preprocessing

1. Drop unnecessary Columns

```
In [36]: # Remove name as it won't help directly for prediction
df.drop(columns=['Name'], inplace=True)
```

2. Drop Missing Values

There are no missing values

3. Encode Categorical Variables

Sex (Male / Female)

```
In [38]: le = LabelEncoder()
df['Sex'] = le.fit_transform(df['Sex'])

# Male = 1, Female = 0
```

4. Feature and target split

```
In [40]: X = df.drop('Survived', axis=1)
y = df['Survived']
```

5. Train-Test Split

```
In [41]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

6. Feature Scaling

```
In [42]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Model Training

Logistic Regression

```
In [43]: model = LogisticRegression()
model.fit(X_train, y_train)
```

Out[43]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

Model Evaluation

Predictions

```
In [45]: y_pred = model.predict(X_test)
```

Accuracy

```
In [46]: accuracy_score(y_test, y_pred)
```

```
Out[46]: 0.7471910112359551
```

Confusion Matrix

```
In [47]: confusion_matrix(y_test, y_pred)
```

```
Out[47]: array([[96, 15],  
                 [30, 37]], dtype=int64)
```

Classification Report

```
In [48]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.76	0.86	0.81	111
1	0.71	0.55	0.62	67
accuracy		0.75	0.75	178
macro avg	0.74	0.71	0.72	178
weighted avg	0.74	0.75	0.74	178

```
In [ ]:
```